




CNRS - Toulouse INP - UT3 - UT Capitole - UT2

Institut de Recherche en Informatique de Toulouse



# Lecture 1: Introduction

***N8EN18B - Contrôle et Apprentissage***

Guilherme IECKER RICARDO

IRIT, Université de Toulouse, CNRS, Toulouse INP, UT3

- 1 Course Organization
- 2 What is this course about?
- 3 The RL Problem
  - Reward
  - Agent and Environment
  - State
  - Inside An RL Agent
  - Problems within RL
- 4 Bibliography



# Course Organization



# Class Info

- Contact: [guilherme.ricardo@irit.fr](mailto:guilherme.ricardo@irit.fr)
- Moodle: Announcements, lecture slides, and additional material
- No official textbook – Recommended reading:
  - An Introduction to Reinforcement Learning, Sutton and Barto (2020).
  - Algorithms for Reinforcement Learning, Szepesvari (2010).
- Assessment:
  - 25% TP
  - 75% Final Exam
- Slides are based on Prof. Silver's course [Silver, 2015]



# Schedule

Class	Date	Time	Topic
CM1	29/03/2024	10h15 - 12h00	Introduction
CM2	02/04/2024	16h15 - 18h00	MDP
CM3	04/04/2024	10h15 - 12h00	Dynamic Programming
CM4	23/04/2024	10h15 - 12h00	
TP	26/04/2024	A: 08h00 - 09h45 R: 10h15 - 12h00	
CM5	14/05/2024	08h00 - 09h45	Deep Q-Learning
Final Exam			

Table: Tentative Schedule




# What is this course about?



# Control and Learning

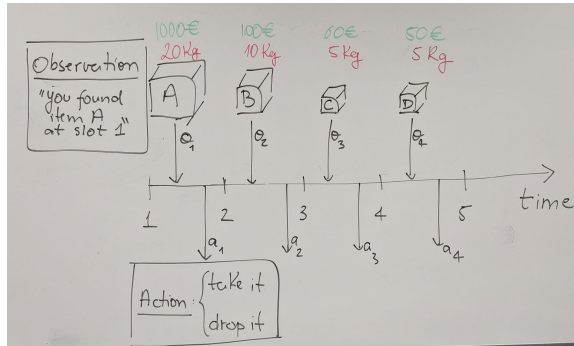
Control = Optimization





# Control and Learning

Learning = Sequential decision making under uncertainty







# Reinforcement Learning




Figure: RL according to different areas.



# Reinforcement Learning




Figure: RL as a Machine Learning paradigm.



# Characteristics of RL

What makes RL different from other ML paradigms?

- There is no supervisor, only a reward signal



# Characteristics of RL

What makes RL different from other ML paradigms?

- There is no supervisor, only a reward signal
- Feedback is delayed, not instantaneous



# Characteristics of RL

What makes RL different from other ML paradigms?

- There is no supervisor, only a reward signal
- Feedback is delayed, not instantaneous
- Time really matters (sequential, non-i.i.d. data)



# Characteristics of RL

What makes RL different from other ML paradigms?

- There is no supervisor, only a reward signal
- Feedback is delayed, not instantaneous
- Time really matters (sequential, non-i.i.d. data)
- Agent's actions affect the subsequent data it receives



# Example of RL Applications

- Fly stunt manoeuvres in a helicopter



# Example of RL Applications

- Fly stunt manoeuvres in a helicopter
- Defeat the world champion at Backgammon





# Example of RL Applications

- Fly stunt manoeuvres in a helicopter
- Defeat the world champion at Backgammon
- Manage an investment portfolio



# Example of RL Applications

- Fly stunt manoeuvres in a helicopter
- Defeat the world champion at Backgammon
- Manage an investment portfolio
- Control a power station



# Example of RL Applications

- Fly stunt manoeuvres in a helicopter
- Defeat the world champion at Backgammon
- Manage an investment portfolio
- Control a power station
- Make a humanoid robot walk



# Example of RL Applications

- Fly stunt manoeuvres in a helicopter
- Defeat the world champion at Backgammon
- Manage an investment portfolio
- Control a power station
- Make a humanoid robot walk
- Play many different video games better than humans



# Example of RL Applications in Networking

5G Drone Relay [Gangula et al., 2018]



# The RL Problem



# The RL Problem

## Reward



- A **reward**  $R_t$  is a scalar feedback signal





# Rewards

- A **reward**  $R_t$  is a scalar feedback signal
- Indicates how well agent is doing at step  $t$



# Rewards

- A **reward**  $R_t$  is a scalar feedback signal
- Indicates how well agent is doing at step  $t$
- The agent's job is to maximize cumulative reward



# Rewards

- A **reward**  $R_t$  is a scalar feedback signal
- Indicates how well agent is doing at step  $t$
- The agent's job is to maximize cumulative reward
- RL is based on the **reward hypothesis**:

## Definition (Reward Hypothesis)

All goals can be described by the maximization of expected cumulative reward.



- A **reward**  $R_t$  is a scalar feedback signal
- Indicates how well agent is doing at step  $t$
- The agent's job is to maximize cumulative reward
- RL is based on the **reward hypothesis**:

### Definition (Reward Hypothesis)

All goals can be described by the maximization of expected cumulative reward.

Do you agree with this statement?



# Examples of Rewards

- Fly stunt manoeuvres in a helicopter
  - +ve reward for following desired trajectory
  - -ve reward for crashing



# Examples of Rewards

- Fly stunt manoeuvres in a helicopter
  - +ve reward for following desired trajectory
  - -ve reward for crashing
- Manage an investment portfolio
  - +ve/-ve reward for each \$ in bank



# Examples of Rewards

- Fly stunt manoeuvres in a helicopter
  - +ve reward for following desired trajectory
  - -ve reward for crashing
- Manage an investment portfolio
  - +ve/-ve reward for each \$ in bank
- Control a power station
  - +ve reward for producing power
  - -ve reward for exceeding safety thresholds



# Examples of Rewards

- Fly stunt manoeuvres in a helicopter
  - +ve reward for following desired trajectory
  - -ve reward for crashing
- Manage an investment portfolio
  - +ve/-ve reward for each \$ in bank
- Control a power station
  - +ve reward for producing power
  - -ve reward for exceeding safety thresholds
- Make a humanoid robot walk
  - +ve reward for forward motion
  - -ve reward for falling over





# Examples of Rewards

- Fly stunt manoeuvres in a helicopter
  - +ve reward for following desired trajectory
  - -ve reward for crashing
- Manage an investment portfolio
  - +ve/-ve reward for each \$ in bank
- Control a power station
  - +ve reward for producing power
  - -ve reward for exceeding safety thresholds
- Make a humanoid robot walk
  - +ve reward for forward motion
  - -ve reward for falling over
- Play many different Atari games better than humans
  - +ve/-ve reward for increasing/decreasing score



# Examples of Rewards

What are the rewards for the 5G relay example?



# Sequential Decision Making

- Goal: select actions to maximize total future reward
- Actions may have long-term consequences
- Reward may be delayed
- It may be better to sacrifice immediate reward to gain more long-term reward
- Examples:
  - A financial investment (may take months to mature)
  - Refuelling a helicopter (might prevent a crash in several hours)
  - Blocking opponent moves (might help winning chances many moves from now)



# The RL Problem

## Agent and Environment





Figure: Agent Abstraction



# Agent and Environment



At each step  $t$ ,

- the agent:
  - Executes action  $A_t$
  - Receives observation  $O_t$
  - Receives scalar reward  $R_t$
- the environment:
  - Receives action  $A_t$
  - Emits observation  $O_{t+1}$
  - Emits scalar reward  $R_{t+1}$
- $t$  increments at environment step

Figure: Agent-Environment cycle



# The RL Problem

## State



# History and State

- The *history* is the sequence of observations, actions, rewards:

$$H_t = O_1, R_1, A_1, \dots, A_{t-1}, O_t, R_t \quad (1)$$





# History and State

- The *history* is the sequence of observations, actions, rewards:

$$H_t = O_1, R_1, A_1, \dots, A_{t-1}, O_t, R_t \quad (1)$$

- i.e., all observable variables up to time  $t$



# History and State

- The *history* is the sequence of observations, actions, rewards:

$$H_t = O_1, R_1, A_1, \dots, A_{t-1}, O_t, R_t \quad (1)$$

- i.e., all observable variables up to time  $t$
- What happens next depends on the history:
  - The agent selects actions
  - The environment selects observation/rewards



# History and State

- The *history* is the sequence of observations, actions, rewards:

$$H_t = O_1, R_1, A_1, \dots, A_{t-1}, O_t, R_t \quad (1)$$

- i.e., all observable variables up to time  $t$
- What happens next depends on the history:
  - The agent selects actions
  - The environment selects observation/rewards
- *State* is the information used to determine what happens next



# History and State

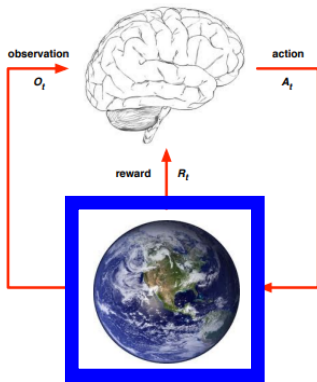
- The *history* is the sequence of observations, actions, rewards:

$$H_t = O_1, R_1, A_1, \dots, A_{t-1}, O_t, R_t \quad (1)$$

- i.e., all observable variables up to time  $t$
- What happens next depends on the history:
  - The agent selects actions
  - The environment selects observation/rewards
- *State* is the information used to determine what happens next
- Formally, the state  $S_t$  at time  $t$  is a function of the history:

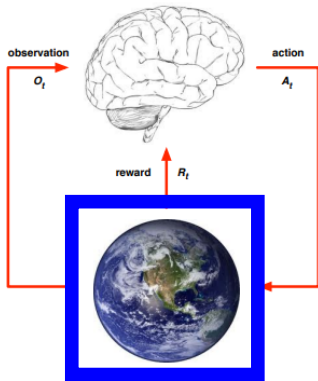
$$S_t = f(H_t) \quad (2)$$

# Environment State



- The **environment state**  $S_t^e$  is the environment's private representation


# Environment State



- The **environment state**  $S_t^e$  is the environment's private representation
- i.e., whatever data the environment uses to pick the next observation/reward



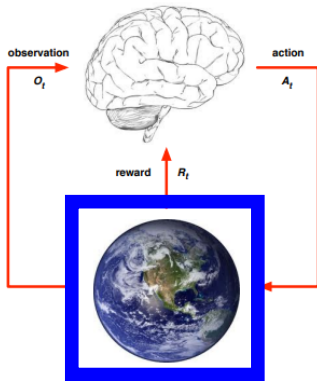
# Environment State



- The **environment state**  $S_t^e$  is the environment's private representation
- i.e., whatever data the environment uses to pick the next observation/reward
- The environment state is not usually visible to the agent




# Environment State



- The **environment state**  $S_t^e$  is the environment's private representation
- i.e., whatever data the environment uses to pick the next observation/reward
- The environment state is not usually visible to the agent
- Even if  $S_t^e$  is visible, it may contain irrelevant information

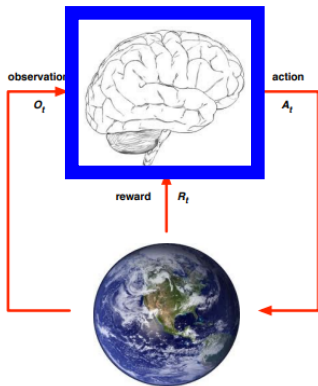


# IRIT Agent State




- The **agent state**  $S_t^a$  is the agent's internal representation

# IRIT Agent State




- The **agent state**  $S_t^a$  is the agent's internal representation
- i.e., whatever information the agent uses to pick the next action

# Agent State



- The **agent state**  $S_t^a$  is the agent's internal representation
- i.e., whatever information the agent uses to pick the next action
- i.e., it is the information used by RL algorithms

# IRIT Agent State



- The **agent state**  $S_t^a$  is the agent's internal representation
- i.e., whatever information the agent uses to pick the next action
- i.e., it is the information used by RL algorithms
- It can be any function of history:

$$S_t^a = f(H_t) \quad (3)$$



# Information State

An information state (a.k.a. Markov state) contains all useful information from the history.



# Information State

An information state (a.k.a. Markov state) contains all useful information from the history.

## Definition

A state  $S_t$  is **Markovian** if, and only if,

$$\mathbb{P}(S_{t+1}|S_t) = \mathbb{P}(S_{t+1}|S_1, \dots, S_t) \quad (4)$$



# Information State

An information state (a.k.a. Markov state) contains all useful information from the history.

## Definition

A state  $S_t$  is **Markovian** if, and only if,

$$\mathbb{P}(S_{t+1}|S_t) = \mathbb{P}(S_{t+1}|S_1, \dots, S_t) \quad (4)$$

- "The future is independent of the past given the present"

$$H_{1:t} \rightarrow S_t \rightarrow H_{t+1:\infty} \quad (5)$$



# Information State

An information state (a.k.a. Markov state) contains all useful information from the history.

## Definition

A state  $S_t$  is **Markovian** if, and only if,

$$\mathbb{P}(S_{t+1}|S_t) = \mathbb{P}(S_{t+1}|S_1, \dots, S_t) \quad (4)$$

- "The future is independent of the past given the present"

$$H_{1:t} \rightarrow S_t \rightarrow H_{t+1:\infty} \quad (5)$$

- Once the state is known, the history may be thrown away





# Information State

An information state (a.k.a. Markov state) contains all useful information from the history.

## Definition

A state  $S_t$  is **Markovian** if, and only if,

$$\mathbb{P}(S_{t+1}|S_t) = \mathbb{P}(S_{t+1}|S_1, \dots, S_t) \quad (4)$$

- "The future is independent of the past given the present"

$$H_{1:t} \rightarrow S_t \rightarrow H_{t+1:\infty} \quad (5)$$

- Once the state is known, the history may be thrown away
- i.e., The state is a sufficient statistic of the future



# Rat Example




Figure: Cognitive Experiment with Rats



# Rat Example




Figure: Cognitive Experiment with Rats

- What if agent state = "last 3 items in sequence"?



# Rat Example




Figure: Cognitive Experiment with Rats

- What if agent state = "last 3 items in sequence"?
- What if agent state = "counts for lights, bells, or levers"?



# Rat Example





Figure: Cognitive Experiment with Rats

- What if agent state = "last 3 items in sequence"?
- What if agent state = "counts for lights, bells, or levers"?
- What if agent state = "complete sequence"?



# Fully Observable Environments



Full observability: agent directly observes environment state

$$O_t = S_t^a = S_t^e \quad (6)$$

- Agent state = environment state = information state
- Formally, this is a **Markov Decision Process (MDP)**

Figure: Agent-Environment cycle



# Partially Observable Environments

- **Partial observability:** agent **indirectly** observes environment:



# Partially Observable Environments

- **Partial observability:** agent **indirectly** observes environment:
  - A robot with a camera vision isn't told its absolute location





# Partially Observable Environments

- **Partial observability:** agent **indirectly** observes environment:
  - A robot with a camera vision isn't told its absolute location
  - A trading agent only observes current prices



# Partially Observable Environments

- **Partial observability:** agent **indirectly** observes environment:
  - A robot with a camera vision isn't told its absolute location
  - A trading agent only observes current prices
  - A poker playing agent only observes public cards



# Partially Observable Environments

- **Partial observability:** agent **indirectly** observes environment:
  - A robot with a camera vision isn't told its absolute location
  - A trading agent only observes current prices
  - A poker playing agent only observes public cards
- Now, agent state  $\neq$  environment state



# Partially Observable Environments

- **Partial observability:** agent **indirectly** observes environment:
  - A robot with a camera vision isn't told its absolute location
  - A trading agent only observes current prices
  - A poker playing agent only observes public cards
- Now, agent state  $\neq$  environment state
- Formally this is a **Partially Observable MDP (POMDP)**



# Partially Observable Environments

- **Partial observability:** agent **indirectly** observes environment:
  - A robot with a camera vision isn't told its absolute location
  - A trading agent only observes current prices
  - A poker playing agent only observes public cards
- Now, agent state  $\neq$  environment state
- Formally this is a **Partially Observable MDP (POMDP)**
- Agent must construct its own state representation  $S_t^a$ , e.g.,



# Partially Observable Environments

- **Partial observability**: agent **indirectly** observes environment:
  - A robot with a camera vision isn't told its absolute location
  - A trading agent only observes current prices
  - A poker playing agent only observes public cards
- Now, agent state  $\neq$  environment state
- Formally this is a **Partially Observable MDP (POMDP)**
- Agent must construct its own state representation  $S_t^a$ , e.g.,
  - Complete history:  $S_t^a = H_t$



# Partially Observable Environments

- **Partial observability**: agent **indirectly** observes environment:
  - A robot with a camera vision isn't told its absolute location
  - A trading agent only observes current prices
  - A poker playing agent only observes public cards
- Now, agent state  $\neq$  environment state
- Formally this is a **Partially Observable MDP (POMDP)**
- Agent must construct its own state representation  $S_t^a$ , e.g.,
  - Complete history:  $S_t^a = H_t$
  - **Beliefs** of environment state:  $S_t^a = (\mathbb{P}(S_t^e = s^1), \dots, \mathbb{P}(S_t^e = s^n))$



# Partially Observable Environments

- **Partial observability:** agent **indirectly** observes environment:
  - A robot with a camera vision isn't told its absolute location
  - A trading agent only observes current prices
  - A poker playing agent only observes public cards
- Now, agent state  $\neq$  environment state
- Formally this is a **Partially Observable MDP (POMDP)**
- Agent must construct its own state representation  $S_t^a$ , e.g.,
  - Complete history:  $S_t^a = H_t$
  - **Beliefs** of environment state:  $S_t^a = (\mathbb{P}(S_t^e = s^1), \dots, \mathbb{P}(S_t^e = s^n))$
  - Recurrent neural network:  $S_t^a = \sigma(S_{t-1}^a W_s + O_t W_o)$





# The RL Problem

## Inside An RL Agent



# Major Components of an RL Agent

An RL agent may include one or more of these components:

- *Policy*: agent's behavior function
- *Value function*: how good is each state and/or action
- *Model*: agent's representation of the environment



- A **policy** is the agent's behavior
- It is a map from state to action, e.g.,
- Deterministic policy:  $a = \pi(s)$
- Stochastic policy:  $\pi(a|s) = \mathbb{P}(A_t = a|S_t = s)$



# Value Function

- **Value function** is a prediction of future reward
- Used to evaluate the goodness/badness of states
- And therefore to select between actions, e.g.,

$$v_{\pi}(s) = \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s] \quad (7)$$




- A **model** predicts what the environment will do next
- Transitions:  $\mathcal{P}$  predicts the next state
- Rewards:  $\mathcal{R}$  predicts the next (immediate) reward, e.g.,

$$\mathcal{P}_{ss'}^a = \mathbb{P}(S_{t+1} = s' | S_t = s, A_t = a) \quad (8)$$

$$\mathcal{R}_{ss'}^a = \mathbb{E}(R_{t+1} | S_t = s, A_t = a) \quad (9)$$



# Maze Example



- Rewards: -1 per time step
- Actions: N, S, E, W
- States: Agent's location

Figure: Maze Example



# Maze Example: Policy and Value Function




Figure: Arrows represent **policy**  $\pi(s)$  for each state  $s$





Figure: Numbers represent **value**  $v_{\pi}(s)$  for each state  $s$



# Maze Example: Model



- Agent may have an internal model of the environment
- Dynamics: how actions change the state
- Rewards: how much reward from each state
- The model may be imperfect

- Grid layout represents transition model  $\mathcal{P}_{ss'}^a$
- Numbers represent immediate reward  $\mathcal{R}_s^a$  from each state  $s$  (same for all  $a$ )





# The RL Problem

Problems within RL



# Exploration and Exploitation

- RL is like trial-and-error learning
- The agent should discover a good policy
- From its experiences of the environment
- Without losing too much reward along the way
- Exploration finds more information about the environment
- Exploitation exploits known information to maximize reward
- **It is usually important to explore as well as exploit**



# Exploration and Exploitation: Examples

- Restaurant Selection  
Exploitation: Go to your favorite restaurant every time  
Exploration: Try a new restaurant
- Online Banner Advertisements  
Exploitation: Show the most successful advertisement  
Exploration: Show a different advertisement
- Oil Drilling  
Exploitation: Drill at the best known location  
Exploration: Drill at a new location
- Game Playing  
Exploitation: Play the move you believe is best  
Exploration: Play an experimental move



# Bibliography



Gangula, R., Esrafilian, O., Gesbert, D., Roux, C., Kaltenberger, F., and Knopp, R. (2018).

Flying rebots: First results on an autonomous uav-based lte relay using open airinterface.

*In 2018 IEEE 19th International Workshop on Signal Processing Advances in Wireless Communications (SPAWC), pages 1–5.*



Silver, D. (2015).

Lectures on reinforcement learning.

url: <https://www.davidsilver.uk/teaching/>.