




CNRS - Toulouse INP - UT3 - UT Capitole - UT2

Institut de Recherche en Informatique de Toulouse



Lecture 5: Temporal-Difference Learning

N8EN18B - Contrôle et Apprentissage

Guilherme IECKER RICARDO

IRIT, Université de Toulouse, CNRS, Toulouse INP, UT3

- 1 On-Policy Control: SARSA
- 2 Off-Policy Control: q -Learning
- 3 Deep Reinforcement Learning
- 4 Final Exam



Recap - Monte-Carlo Learning

Access the Python Notebook:

<https://guilhermeir.github.io/teaching/rl/mc.ipynb>



Recap - TD Learning: Prediction

Tabular TD(0) for estimating v_π

Input: the policy π to be evaluated

Algorithm parameter: step size $\alpha \in (0, 1]$

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

$A \leftarrow$ action given by π for S

 Take action A , observe R, S'

$V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$

$S \leftarrow S'$

 until S is terminal



Recap - Bellman Equations

■ Bellman Expectation Equation

$$\begin{aligned} v_{\pi}(s) &= \sum_{a \in \mathcal{A}} \pi(a|s) \cdot \left(\mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \cdot v_{\pi}(s') \right) \\ q_{\pi}(s, a) &= \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \sum_{a' \in \mathcal{A}} \pi(a'|s') \cdot q_{\pi}(s', a') \end{aligned} \quad (1)$$

■ Bellman Optimality Equation

$$\begin{aligned} v_*(s) &= \max_{a \in \mathcal{A}} \left\{ \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \cdot v_*(s') \right\} \\ q_*(s, a) &= \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \cdot \max_{a' \in \mathcal{A}} \{ q_*(s', a') \} \end{aligned} \quad (2)$$



On-Policy Control: SARSA




MC vs. TD Control

- Temporal-Difference (TD) Learning has several advantages over Monte-Carlo (MC) Learning
 - Lower variance
 - Online
 - Incomplete sequences
- Natural idea: use TD instead of MC in our control loop (value iteration), i.e.,
 - Apply TD to $Q(S, A)$
 - Use ϵ -Greedy policy improvement
 - Update every time-step




Updating Action-Value Functions with SARSA



$$Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma Q(S', A') - Q(S, A))$$



On-Policy Control with SARSA



Every time-step:

Policy evaluation: SARSA, $Q \approx q_\pi$

Policy improvement: ϵ -Greedy Policy Improvement



SARSA Algorithm for On-Policy Control

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Loop for each step of episode:

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

 until S is terminal



Off-Policy Control: q -Learning



Off-Policy Learning

- Evaluate target policy $\pi(a|s)$ to compute $v_\pi(s)$ or $q_\pi(s, a)$
- While following behavior policy $\mu(a|s)$

$$\{S_1, A_1, R_2, \dots, S_T\} \sim \mu$$

- Why is this important?
 - Learn from observing humans or other agents
 - Re-use experience generated from old policies $\pi_1, \pi_2, \dots, \pi_{t-1}$
 - Learn about optimal policy while following exploratory policy
 - Learn about multiple policies while following one policy



Q-Learning

- We now consider off-policy learning of action-values $Q(s, a)$
- Next action is chosen using behavior policy $A_{t+1} \sim \mu(\cdot|S_t)$
- But we consider alternative successor action $A' \sim \pi(\cdot|S_t)$
- And update $Q(S_t, A_t)$ towards value of alternative action

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A') - Q(S_t, A_t))$$



Off-Policy Control with Q-Learning

- We now allow both behavior and target policies to **improve**
- The target policy π is **greedy** w.r.t. $Q(s, a)$


$$S(S_{t+1} = \arg \max_{a'} Q(S_{t+1}, a') \quad (3)$$

- The behavior policy μ is, e.g., ϵ -**greedy** w.r.t. $Q(s, a)$
- The Q-Learning target then simplifies:

$$\begin{aligned} & R_{t+1} + \gamma Q(S_{t+1}, A') \\ &= R_{t+1} + \gamma Q(S_{t+1}, \arg \max_{a'} Q(S_{t+1}, a')) \\ &= R_{t+1} + \max_{a'} \gamma Q(S_{t+1}, a') \end{aligned} \quad (4)$$



Q-Learning Control Algorithm



$$Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma \max_{a'} Q(S', a') - Q(S, A))$$

Theorem

Q-Learning control converges to the optimal action-value function,
 $Q(s, a) \rightarrow q_*(s, a)$



Q-Learning Algorithm for Off-Policy Control

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal



Deep Reinforcement Learning



Deep Reinforcement Learning

Additional set of slides



Final Exam



How to study?

- Class slides (available [here](#))
- Reading (text-book available [here](#))
 - Chapter 1: all sections
 - Chapter 3: all sections
 - Chapter 4: all sections except 4.5
 - Chapter 5: sections 5.1 - 5.4
 - Chapter 6: sections 6.1 - 6.5
- Studying examples and solving problems from text-book (included in the sections above)
 - [Here](#) you have the code for all examples in the book
 - [Here](#) you have all solutions for the book's questions
- This [list](#) of exercises.



Rules

- 1-hour long final exam
- There are 20 points + 5 bonus points, i.e., your **actual grade** $\in [0, 25]$
- Your **final grade** is $\min(20, \text{actual grade})$
- Only 1 page (1 side of a sheet of paper) of personal notes is allowed
- No devices, books, etc.
- Zero cheating tolerance!